

# Analysis of Multi-Join Benchmarking

Harshad Deshmukh  
University of Wisconsin-Madison  
harshad@cs.wisc.edu

Jignesh M. Patel  
University of Wisconsin-Madison  
jignesh@cs.wisc.edu

## ABSTRACT

### 1. INTRODUCTION

Performance of join operations is crucial for relational data analytics. For a long time, database community has tried to improve the performance of join operations. The two main approaches in join algorithm designs are - partition-based and no partition-based join implementations.


Remarkably, most of the prior work on join performance [1, 2] has focused on a single join, usually implemented with synthetic schema and data. The prior work is valuable in understanding the implications of various parameters on the performance of one join operation. However, real life decision support queries often involve multiple join operations (multi-join), typically executed one after another on a dataset organized as star-schema or snowflake-schema. The debate on partitioning-based *vs* no partitioning-based joins has not been extended to the settings of multi-joins.

An additional complexity associated with analyzing multi-joins is the potential re-partitioning operation between two joins. The re-partitioning is needed when the input to the partitioned join is either not partitioned at all or is partitioned on a different attribute than the join attribute.

Our goal is to devise a high performance hash-based multi-join implementation algorithm for a right deep pipeline.  $\rightarrow$  (Justify why we choose only hash-based implementations)  $\leftarrow$ . One example of a right deep pipeline is shown in Figure 1.

An important aspect of the right-deep pipeline in Figure 1 is that the build side in all the join operations is always some dimension table. We will leverage this aspect of the right-deep pipeline in our algorithm design.  $\rightarrow$  (Think of extending the idea to left-deep or zig-zag trees.)  $\leftarrow$

Also note that Figure 1 is a *logical plan* produced by the optimizer. It gives information about join ordering. e.g. join between *Fact* and *Dim<sub>1</sub>* relations should be performed first, followed by the join of the first join's result and *Dim<sub>2</sub>* relation. Figure 1 does not specify the type of join (e.g. hash-based or sort-merge) and/or the partitioning nature of



**Figure 1: A right-deep pipeline consisting of three join operations.  $Dim_i$  indicates  $i^{th}$  dimension table and *Fact* denotes the fact table in a star-schema.**

the join such as radix-based partitioned join or no partition-based join.

The goal of our algorithm is to produce a high performance *physical plan* which consists of the partitioning nature of each join in the pipeline. We formally describe the problem in the next section.

### 1.1 Problem Statement

Given a right deep pipeline (as shown in Figure 1), consisting of  $k$  join operations viz.  $J_1, J_2, \dots, J_k$ , between fact and dimension tables, find the partitioning characteristic of the implementation of each join  $p(J_i)$ , which can result in overall high performance of the pipeline.

Our objective is to minimize the execution time of the multi-join pipeline. Let us define the execution time of join operation (probe + build phase)  $J_i$  as  $T(J_i)$ , or simply  $T_i$ .

Therefore the objective is to minimize  $\sum_{i=1}^k T_i$ .

Note that if the execution exhibits pipelined parallelism, the above objective might be incorrect. In such a case, the total execution time for the pipeline is simply the time difference between end of execution of the last join and beginning of execution of the first join.

We consider only two options for the partitioning characteristic of the join implementation, i.e.  $p(J_i) := \text{PARTITION-BASED}$  or  $p(J_i) := \text{NO PARTITION-BASED}$ . Thus, our problem can be viewed as an assignment problem, where we have to assign a value to  $k$  binary variables. Therefore, the size of the solution space is  $2^k$ , i.e. there are  $2^k$  number

of different possible physical plans.

In the next section, we discuss some proposals to solve the above problem.

## 1.2 Proposed Methodology

We first discuss the parameters involved in this problem. The execution time  $T_i$  per join operation  $J_i$  consists of both build and probe phases. If the join implementation is partitioning-based, then the time to partition (or re-partition) the relations is included in  $T_i$ .

We use a metric to differentiate the performances of partitioned and non-partitioned probes. This metric is the number of CPU cycles per probing tuple ( $C_{probe}$ ), and it is impacted by the size of the probed hash table  $Size_{HT}$ . In the two implementations, viz. partitioned and non-partitioned the relationship between  $C_{probe}$  and  $Size_{HT}$  is different. We study this relationship and estimate the cost of a given probe operation. In the next section, we present some initial empirical results to understand this relationship.

The partitioned implementation of a join operator may consist of partitioning cost. Therefore, we define another metric  $C_{part}$ , which is the number of CPU cycles spent in partitioning a single input tuple.

Finally, we estimate the cost of each join operations in terms of CPU cycles per input tuple.

We may use a well known technique such as dynamic programming to solve the assignment problem.

Note that we restrict ourselves to one form of materialization of the results i.e. eager materialization. This means that each join operation's result is immediately materialized as soon as the operation is complete. (An alternative could be late materialization, as is used by MonetDB.) *→ (Unanswered question: Do assignment change as the query is in progress? e.g. If there are skewed results of the lower join operations in the query plan.) ←*

### 1.2.1 Heuristics

As described earlier, the total number of possible physical plans is exponential in the number of join operations. It may be prohibitively expensive to analyze all possible plans. Therefore, we can prune the number of joins by certain heuristics. One such heuristic could be: flip-flop between partitioned and non-partitioned joins is not desirable for successive join operations.

## 2. EXPERIMENTAL EVALUATION

Evaluation description goes here.

## 3. CONCLUSIONS

Conclusions go here.

## 4. RELATED WORK

The database community has paid significant attention to improve join performance in the past decade. Blanas et al. [1] analyzed the join performance in single socket, large main-memory server environments using various join implementations such as radix-based partitioning, no-partitioned join and using uniform as well as skewed data sets.

Schuh et al. [2] reviewed various join implementations from the research literature from the past decade. They classified the join algorithms in two categories - partitioning

based and no-partitioning based join algorithms, and compared their performance characteristics. As they observed, prior work on join performance has mostly focused on synthetic queries and synthetic datasets. In the interest of discussing the join algorithm performance in the context of real life queries, they measured the performance of various join algorithms on a real query (TPC-H Query 19).

## 5. REFERENCES

- [1] S. Blanas, Y. Li, and J. M. Patel. Design and evaluation of main memory hash join algorithms for multi-core cpus. In *SIGMOD*, pages 37–48, 2011.
- [2] S. Schuh, X. Chen, and J. Dittrich. An experimental comparison of thirteen relational equi-joins in main memory. In *SIGMOD*, pages 1961–1976, 2016.